



Introduction to OpenMP

First Part

Silvio Stanzani , Raphael Cóbe , Rogério Iope, Jefferson Fialho
UNESP – Center for Scientific Computing

silvio@ncc.unesp.br , rmcobe@ncc.unesp.br , rogerio@ncc.unesp.br
, jfialho@ncc.unesp.br

UNESP Center for Scientific Computing

- Consolidates scientific computing resources for São Paulo State University (UNESP) researchers
 - It mainly uses Grid computing paradigm
- Main users
 - UNESP researchers, students, and software developers
 - SPRACE (São Paulo Research and Analysis Center) physicists and students
 - ❑ Caltech, Fermilab, CERN
 - ❑ São Paulo CMS Tier-2 Facility
- P&D Projects
 - Intel and Huawei

Recommendations

- Desired Requirements
 - Basic knowledge about C programming
 - Basic knowledge about how to write algorithms
 - Knowledge about how to run programs in Torque
- Do not let the class continue if you have doubt!
 - Please INTERRUPT if you have doubt (english or spanish or portuguese)

Agenda

- Theory 1:30
 - Parallel processing
 - Process / Thread
 - OpenMP Programming Model
 - ❑ Run-Time Library Routines
 - ❑ Environment Variables
 - ❑ Directives
 - Parallel region
 - Parallel loop
 - Data Environment
 - Coffee
 - Hands-on 2:00
 - Lunch
- Theory 1:30
 - Race conditions
 - Synchronization Constructs
 - Race condition detection
 - Task Parallelism
 - Performance evaluation with Valgrind
 - coffee
 - Hands-on 2:00
 - Overtime

Computer Architecture

- Uniprocessor architecture
 - New processor generations increased performance (Moore's Law)
 - Until processor performance achieved its limit
- Trends to keep increasing performance:
 - More processing units in a single processor
 - More processors in the same machine
- In order to achieve better performance, developers **MUST** explore parallelism.

Parallel Processing

- ❑ A parallel computer is a computer system that uses multiple processing elements simultaneously in a cooperative manner to solve a computational problem
- ❑ Parallel processing includes techniques and technologies that make it possible to compute in parallel
 - Hardware, networks, operating systems, parallel libraries, languages, compilers, algorithms, tools, ...
- ❑ Parallel computing is an evolution of serial computing
 - Parallelism is natural
 - Computing problems differ in level / type of parallelism

Shared Memory Architecture

- Complex Memory System Architecture
- Transparent to Users
- Influences performance!

Main Memory

Multi Level Cache

Processing
Unit 1

Processing
Unit 2

...


Processing
Unit N

Concurrency and Synchronization

- Consider a program composed of multiple tasks to be executed in a computer
 - Tasks are **concurrent** with they can execute at the same time (concurrent execution)
- If a task requires results produced by other tasks in order to execute correctly, the task's execution is dependent
 - Some form of **synchronization** must be used to enforce (satisfy) dependencies

Concurrency and Synchronization

— Example:



- Global Variables
- I/O Operations

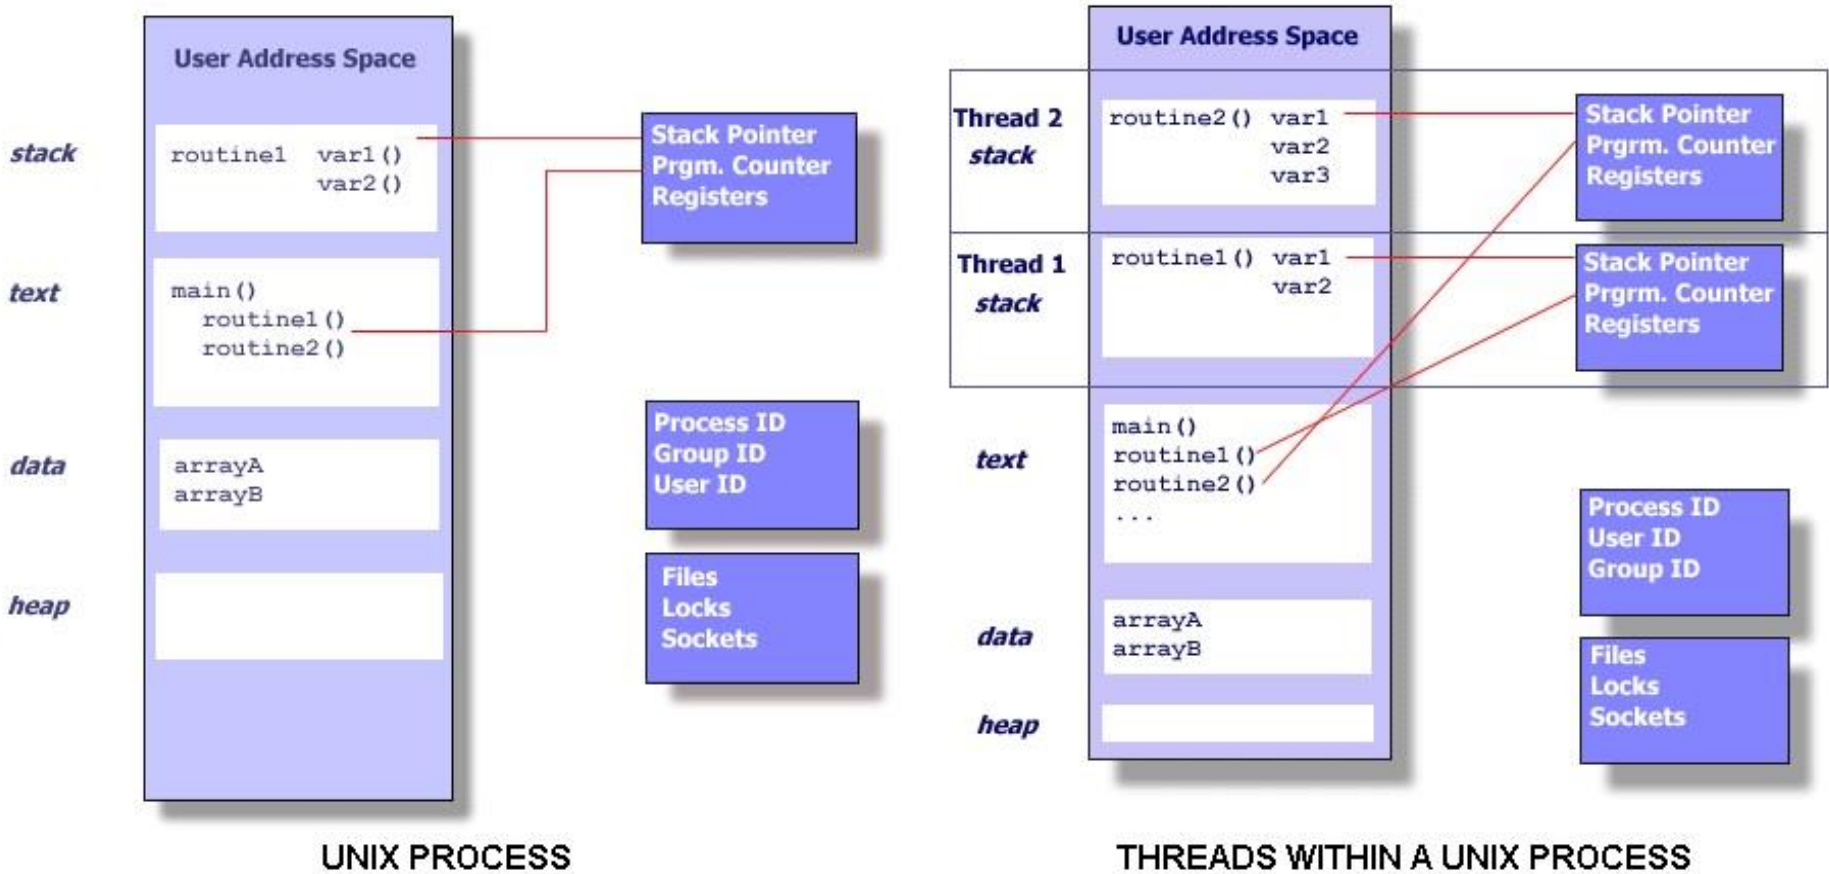
Parallelism

- Parallelism = concurrency + “parallel” hardware
 - Find concurrent execution opportunities
 - Develop application to execute in parallel
 - Run application on parallel hardware
- Motivations for parallelism
 - Faster time to solution (response time)
 - Solve bigger computing problems (in same time)
 - Effective use of machine resources

Process and Threads

- Any program running on Operational system is called Process:
 - Is composed of at least of one Thread
 - Can fork several threads which is a copy of itself
- Creating a new thread are much faster than create a new process
- There are libraries that support thread creation such as Pthreads

Process and Threads



Source: <https://computing.llnl.gov/tutorials/pthreads/>

Multithreaded Programming

- **Multithreading** is the ability of a O.S. to execute one process using several resources simultaneously by the means of threads
- **Multithreaded Programming** is a parallel programming technique that has the objective of prepare your program to be executed as concurrent parts on several threads
- **Pthread** is one library for Multithreaded Programming

Pthreads Example

```
#include <pthread.h>
```

```
void *inc_x(void *x_void_ptr){  
    int *x_ptr = (int *)x_void_ptr;  
    while(++(*x_ptr) < 100);  
    printf("x increment finished\n");  
  
    return NULL;  
}
```

```
int main(){
```

```
    int x = 0, y = 0;  
    pthread_t inc_x_thread;
```

```
    printf("x: %d, y: %d\n", x, y);
```

```
    pthread_create(&inc_x_thread, NULL, inc_x,  
    &x)
```

```
    while(++y < 100);
```

```
    printf("y increment  
finished\n");
```

```
    pthread_join(inc_x_thread,  
    NULL)
```

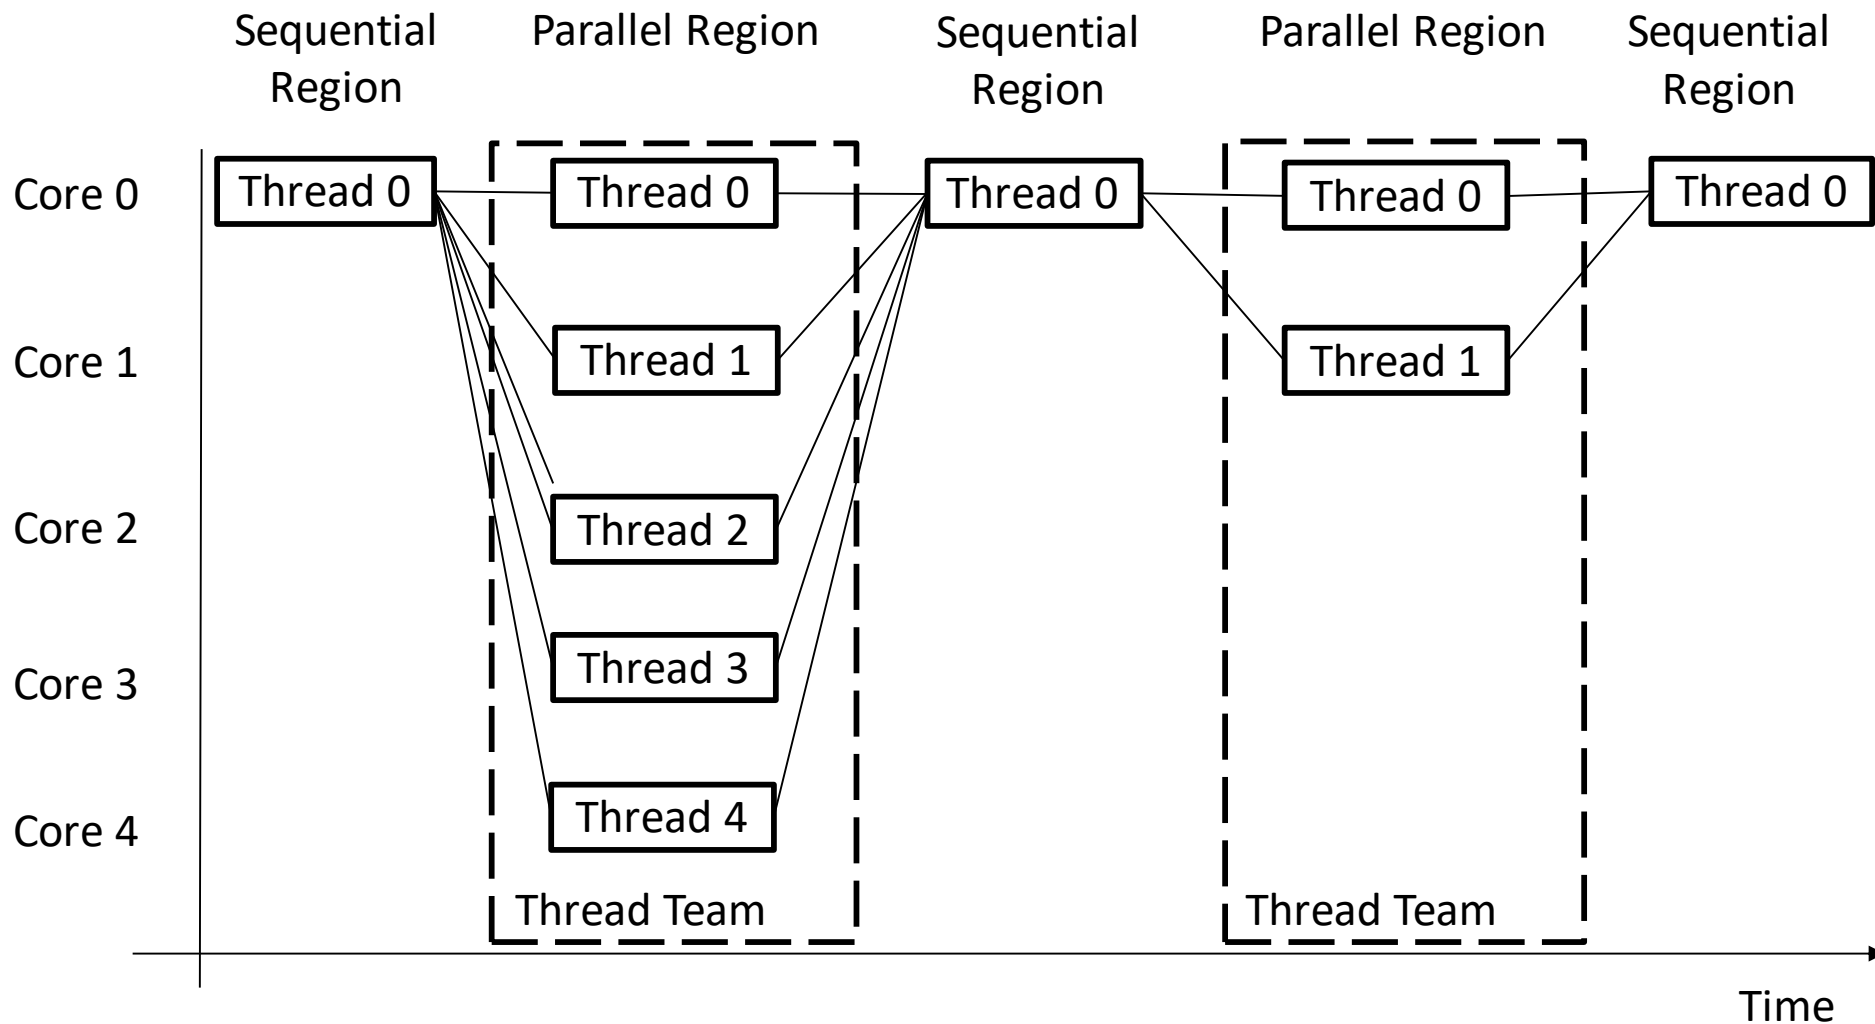
```
    printf("x: %d, y: %d\n", x, y);
```

```
    return 0;  
}
```

OpenMP

- OpenMP is an acronym for Open Multi-Processing
- An Application Programming Interface (API) for developing parallel programs in shared memory architectures
 - API based on Pragmas – C code extensions
- Three primary components of the API are:
 - Compiler Directives
 - Runtime Library Routines
 - Environment Variables
- De facto standard - specified for C / C++ and FORTRAN
- <http://www.openmp.org/>
 - Specification, examples, tutorials and documentation

OpenMP



OpenMP - Core elements

Parallel control structures

Form a team of threads and execute them in parallel

```
omp parallel
```

Synchronization

Coordinates thread execution

```
omp atomic
omp barrier
omp critical
omp flush
omp master
omp ordered
omp taskgroup
omp taskwait
```

Work sharing

Distribute work among threads

```
omp [parallel] loop
omp [parallel] sections
omp [parallel] workshare
omp single
```

Data environment

Control variables scope

```
omp threadprivate
shared/*private
clauses
```

OpenMP - Core elements

OpenMP 4.0 - Co-Processors and Accelerators

SIMD vectorization

Offload execution

Thread affinity

Tasking

Structures for deferring execution

```
omp task  
omp taskyield
```

Runtime environment

Runtime functions and environment
variables

```
omp_set_num_threads(), etc.  
OMP_SCHEDULE, etc.
```

Loop

- Serial Application example:

```
Int i=0;
```

```
N=25;
```

```
for (i=0; i<N; i++)
```

```
    a[i] = a[i] + b;
```

- Iterations of a loop represents tasks that can be executed concurrently;

Parallel Region

```
#pragma omp parallel
```

```
{
```

```
... //Code that need to be executed concurrently goes here
```

```
}
```

- The region enclosed by **pragma omp parallel** will be execute by all threads
- Loop iterations can be divided among threads

OpenMP Sample Program

```
#include <stdio.h>
```

```
int main() {  
    char hn[600];
```

```
    #pragma omp parallel
```

```
{
```

```
    gethostname(hn,600);
```

```
    printf("hello from hostname %s %d\n",hn);
```

```
}
```

```
    return(0);
```

```
}
```

Compiling and running an OpenMP application

- Build the application using gcc

```
gcc <source-code> -o <omp_binary> -fopenmp
```

- Build the application using pgi

```
pgcc <source-code> -o <omp_binary> -mp
```

- Launch the application

```
export OMP_NUM_THREADS=10
```

```
./omp_binary
```

Executing on Los Andes Cluster

- Create job file:

- Ex: job1

```
#!/bin/bash
```

```
#PBS -k o
```

```
#PBS -l nodes=1:ppn=1,walltime=30:00
```

```
#PBS -M jthutt@tatooine.net
```

```
#PBS -m abe
```

```
#PBS -N FirstOMPjob
```

```
#PBS -j oe
```

```
/hpcfs/home/sshpc/jobs/execjob.sh
```

- Submit job qsub

- Qsub job1

- Verify job status qstat

- qstat

OpenMP Functions

- `omp_get_max_threads()`
 - Amount of processing units (cores)
- `omp_get_thread_limit()`
 - Amount of threads that O.S. can Manage
- `omp_get_thread_num();`
 - Get the thread id
- `omp_set_num_threads(8);`
 - Setup the amount of threads to be used
- Environmental variables:
 - `OMP_NUM_THREADS`: define the amount of threads to execute a program using OpenMP
 - ❑ Example: `export OMP_NUM_THREADS=10`

Data Environment

- How threads communicates?
 - Using variables (global and local)
- OpenMP Allows developers to define variables to be private or global among other(Attribute Clauses):
 - `shared(list)` : global variable accross all threads
 - `private(list)`: each thread has its own version, initial value is 0
 - `firstprivate(list)`: each thread has its own version, initial value is the last version before OpenMP region

Loop sharing

- Create a parallel region and distribute the loop iterations among the threads
 - Complete version:
 - ❑ `Pragma omp parallel`
 - `Pragma omp for`
 - short version:
 - ❑ `Pragma omp parallel for`
- `lastprivate(list)`: each thread has its own version, after the end of openmp region the variable receives the value from last thread

Hands-on

- option-price
- Matrix Transposition
- IronBar (ERAD-SP 2018 warmup)



Introduction to OpenMP

Second Part

Silvio Stanzani , Raphael Cóbe , Rogério Iope, Jefferson Fialho
UNESP – Center for Scientific Computing

silvio@ncc.unesp.br , rmcobe@ncc.unesp.br , rogerio@ncc.unesp.br
 , jfialho@ncc.unesp.br

Agenda

- Theory 1:30
 - Parallel processing
 - Process / Thread
 - OpenMP Programming Model
 - Run-Time Library Routines
 - Environment Variables
 - Directives
 - ❑ Parallel region
 - ❑ Parallel loop
 - ❑ Data Environment
 - Coffee
 - Hands-on 2:00
 - lunch
- Theory 1:30
 - Race condition
 - Synchronization Constructs
 - Performance evaluation with Valgrind
 - Race condition detection
 - Task parallelism
 - coffee
 - Hands-on 2:00
 - Overtime

Race Condition

- When two or more threads perform operations on shared data, it is impossible to know the order in which this operations will be performed;
- This is a condition in which one or more threads are "racing" to perform the same operation
- The program will not end with a bug, but in some cases will return with incorrect results

Race Condition

- Example of a race condition:

```
#pragma omp parallel for
for(i=0 ; i<size_of_input_array; i++)
{
    Int *tmpsum = input+i;
    sum += *tmpsum;
}
```

- Every execution return different results!

Race Condition

- Solution to solve race condition problems:
 - Break the dependency using a different algorithm
 - Enforce synchronization: the execution is performed sequentially by all threads
- OpenMP provides several options for synchronization
- Synchronization enforce performance penalties!

Breaking Dependencies

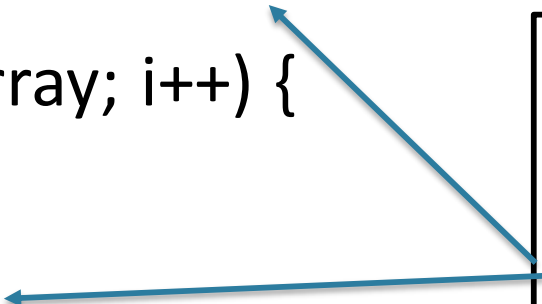
- Example...
- Algorithm from first challenge;

Synchronization

- Synchronization directives:
 - `omp atomic`:
 - ❑ Ensures that a specific memory location is updated atomically, which prevents the possibility of multiple, simultaneous reading and writing of threads.
 - `omp critical`
 - ❑ Specifies a code block that is restricted to access by only one thread at a time.
 - `omp ordered`
 - ❑ Specifies a code block in a worksharing loop that will be run in the *order* of the loop iterations

Synchronization


```
#pragma omp parallel for ordered  
for(i=0 ; i<size_of_input_array; i++) {  
    Int *tmpsum = input+i;  
    #pragma omp ordered  
    sum += *tmpsum;  
}
```



You must put
ordered
clause in a
loop with
ordered
clause

Synchronization


```
#pragma omp parallel for  
for(i=0 ; i<size_of_input_array; i++) {  
    Int *tmpsum = input+i;  
    #pragma omp critical  
    {  
        sum += *tmpsum;  
    }  
}
```



You can enclose
a code region
inside critical
clause

Synchronization

```
#pragma omp parallel for ordered  
for(i=0 ; i<size_of_input_array; i++) {  
    Int *tmpsum = input+i;  
    #pragma omp atomic  
    sum += *tmpsum;  
}
```



Atomic can
embrace a
single line only

Synchronization

- Synchronization directives:
 - `omp barrier`
 - ❑ Specifies a point in the code where each thread must wait until all threads in the team arrive.
 - `omp master`
 - ❑ Specifies the beginning of a code block that must be executed only once by the master thread of the team.
 - `omp single`
 - ❑ Only one thread execute the code block

Synchronization

- Barrier, master and single example

Detect Race Condition


- Detect race condition can be difficult!
- How to detect race condition?
 - Manual Debugging
 - profile tools: ex: valgrind
- Valgrind:
 - Compile using flag -g
 - ❑ Example: `g++ -g -fopenmp sum.cpp -o sum`
 - Execute with helgrind
 - ❑ `valgrind --tool=helgrind ./sum < sum.in &> t`

Task Parallelism

- Tasks are independent units of work
 - code to execute
 - Input/output data
- Threads are assigned to perform the work of each task.
- Tasks can be defined as a relation of dependency

Task Parallelism

- Task Parallelism model of OpenMP.



Task Parallelism

- tasks must be created inside of a parallel region:
 - **#pragma omp task**

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        printf("hello world\n");

        #pragma omp task
        printf("hello again!\n");
    }
}
```

Task Parallelism

```
int fib (int n) {  
    int x,y;  
    if ( n < 2 ) return n;  
  
    x = fib(n-1);  
  
    y = fib(n-2);  
  
    return x+y;  
}  
  
int main() {  
    int NN = 46;  
  
    fib(NN);  
}
```

```
int fib (int n) {  
    int x,y;  
    if ( n < 2 ) return n;  
    #pragma omp task shared(x) if(n>30)  
    x = fib(n-1);  
    #pragma omp task shared(y) if(n>30)  
    y = fib(n-2);  
    #pragma omp taskwait  
    return x+y;  
}  
  
int main() {  
    int NN = 460;  
    #pragma omp parallel  
    {  
        #pragma omp master  
        fib(NN);  
    }  
}
```

Hands-on

- Harmonic Progression SUM (WSCAD 2016 warmup)
 - Best place for synchronization
- Apply task parallelism to QuickSort algorithm

Harmonic progression sum

- `g++ -O3 -g -fopenmp sum.cpp -o sum`

Problem A

Harmonic progression sum

The simplest harmonic progression is

$$1/2, 1/3, 1/4, 1/5, \dots$$

Let $S_n = \sum_{i=1}^n (1/i)$, compute this sum to arbitrary precision after the decimal point.

Input

The input contains only one test case. The first line contains two values; the first is the number of digits D and the second is the value of N . Consider $(1 \leq D \leq 10^7)$ and $(1 \leq N \leq 10^8)$.

The input must be read from the standard input.

Output

The output contains only one line printing the value of the sum with exact D precision.

The output must be written to the standard output.

Reference

- Structured parallel programming
- Shared memory application programming